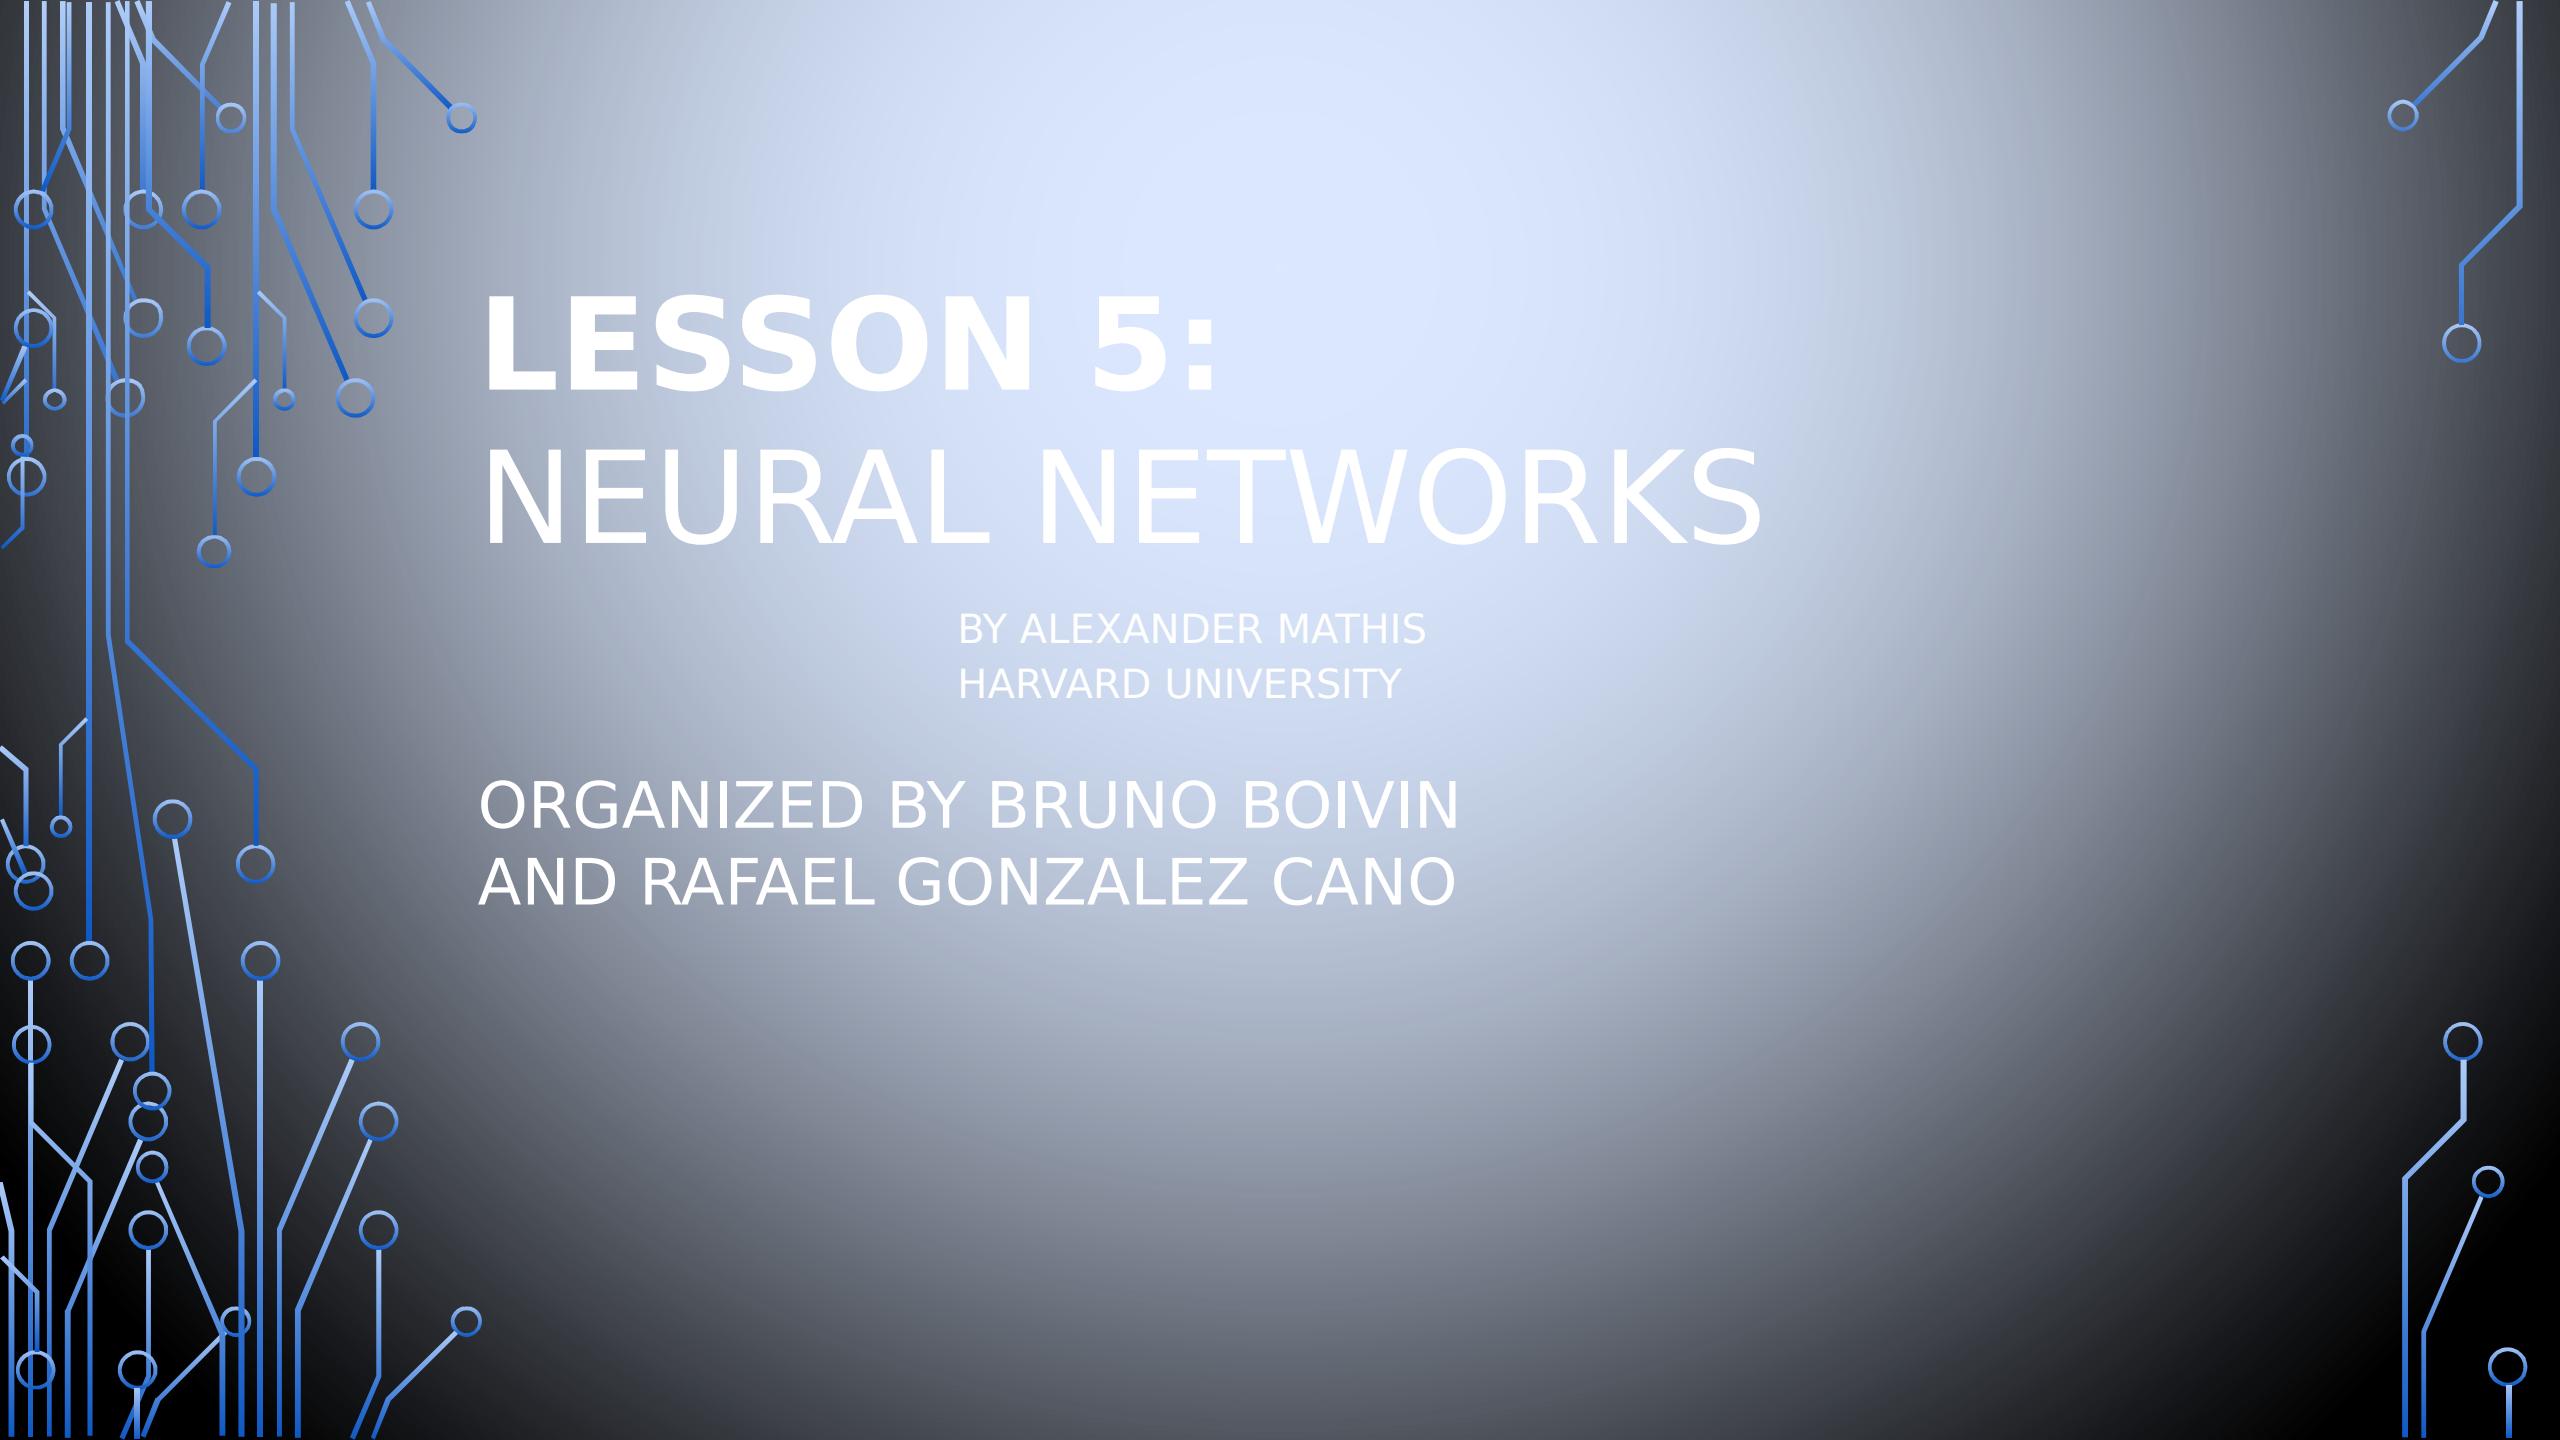




neuro dev

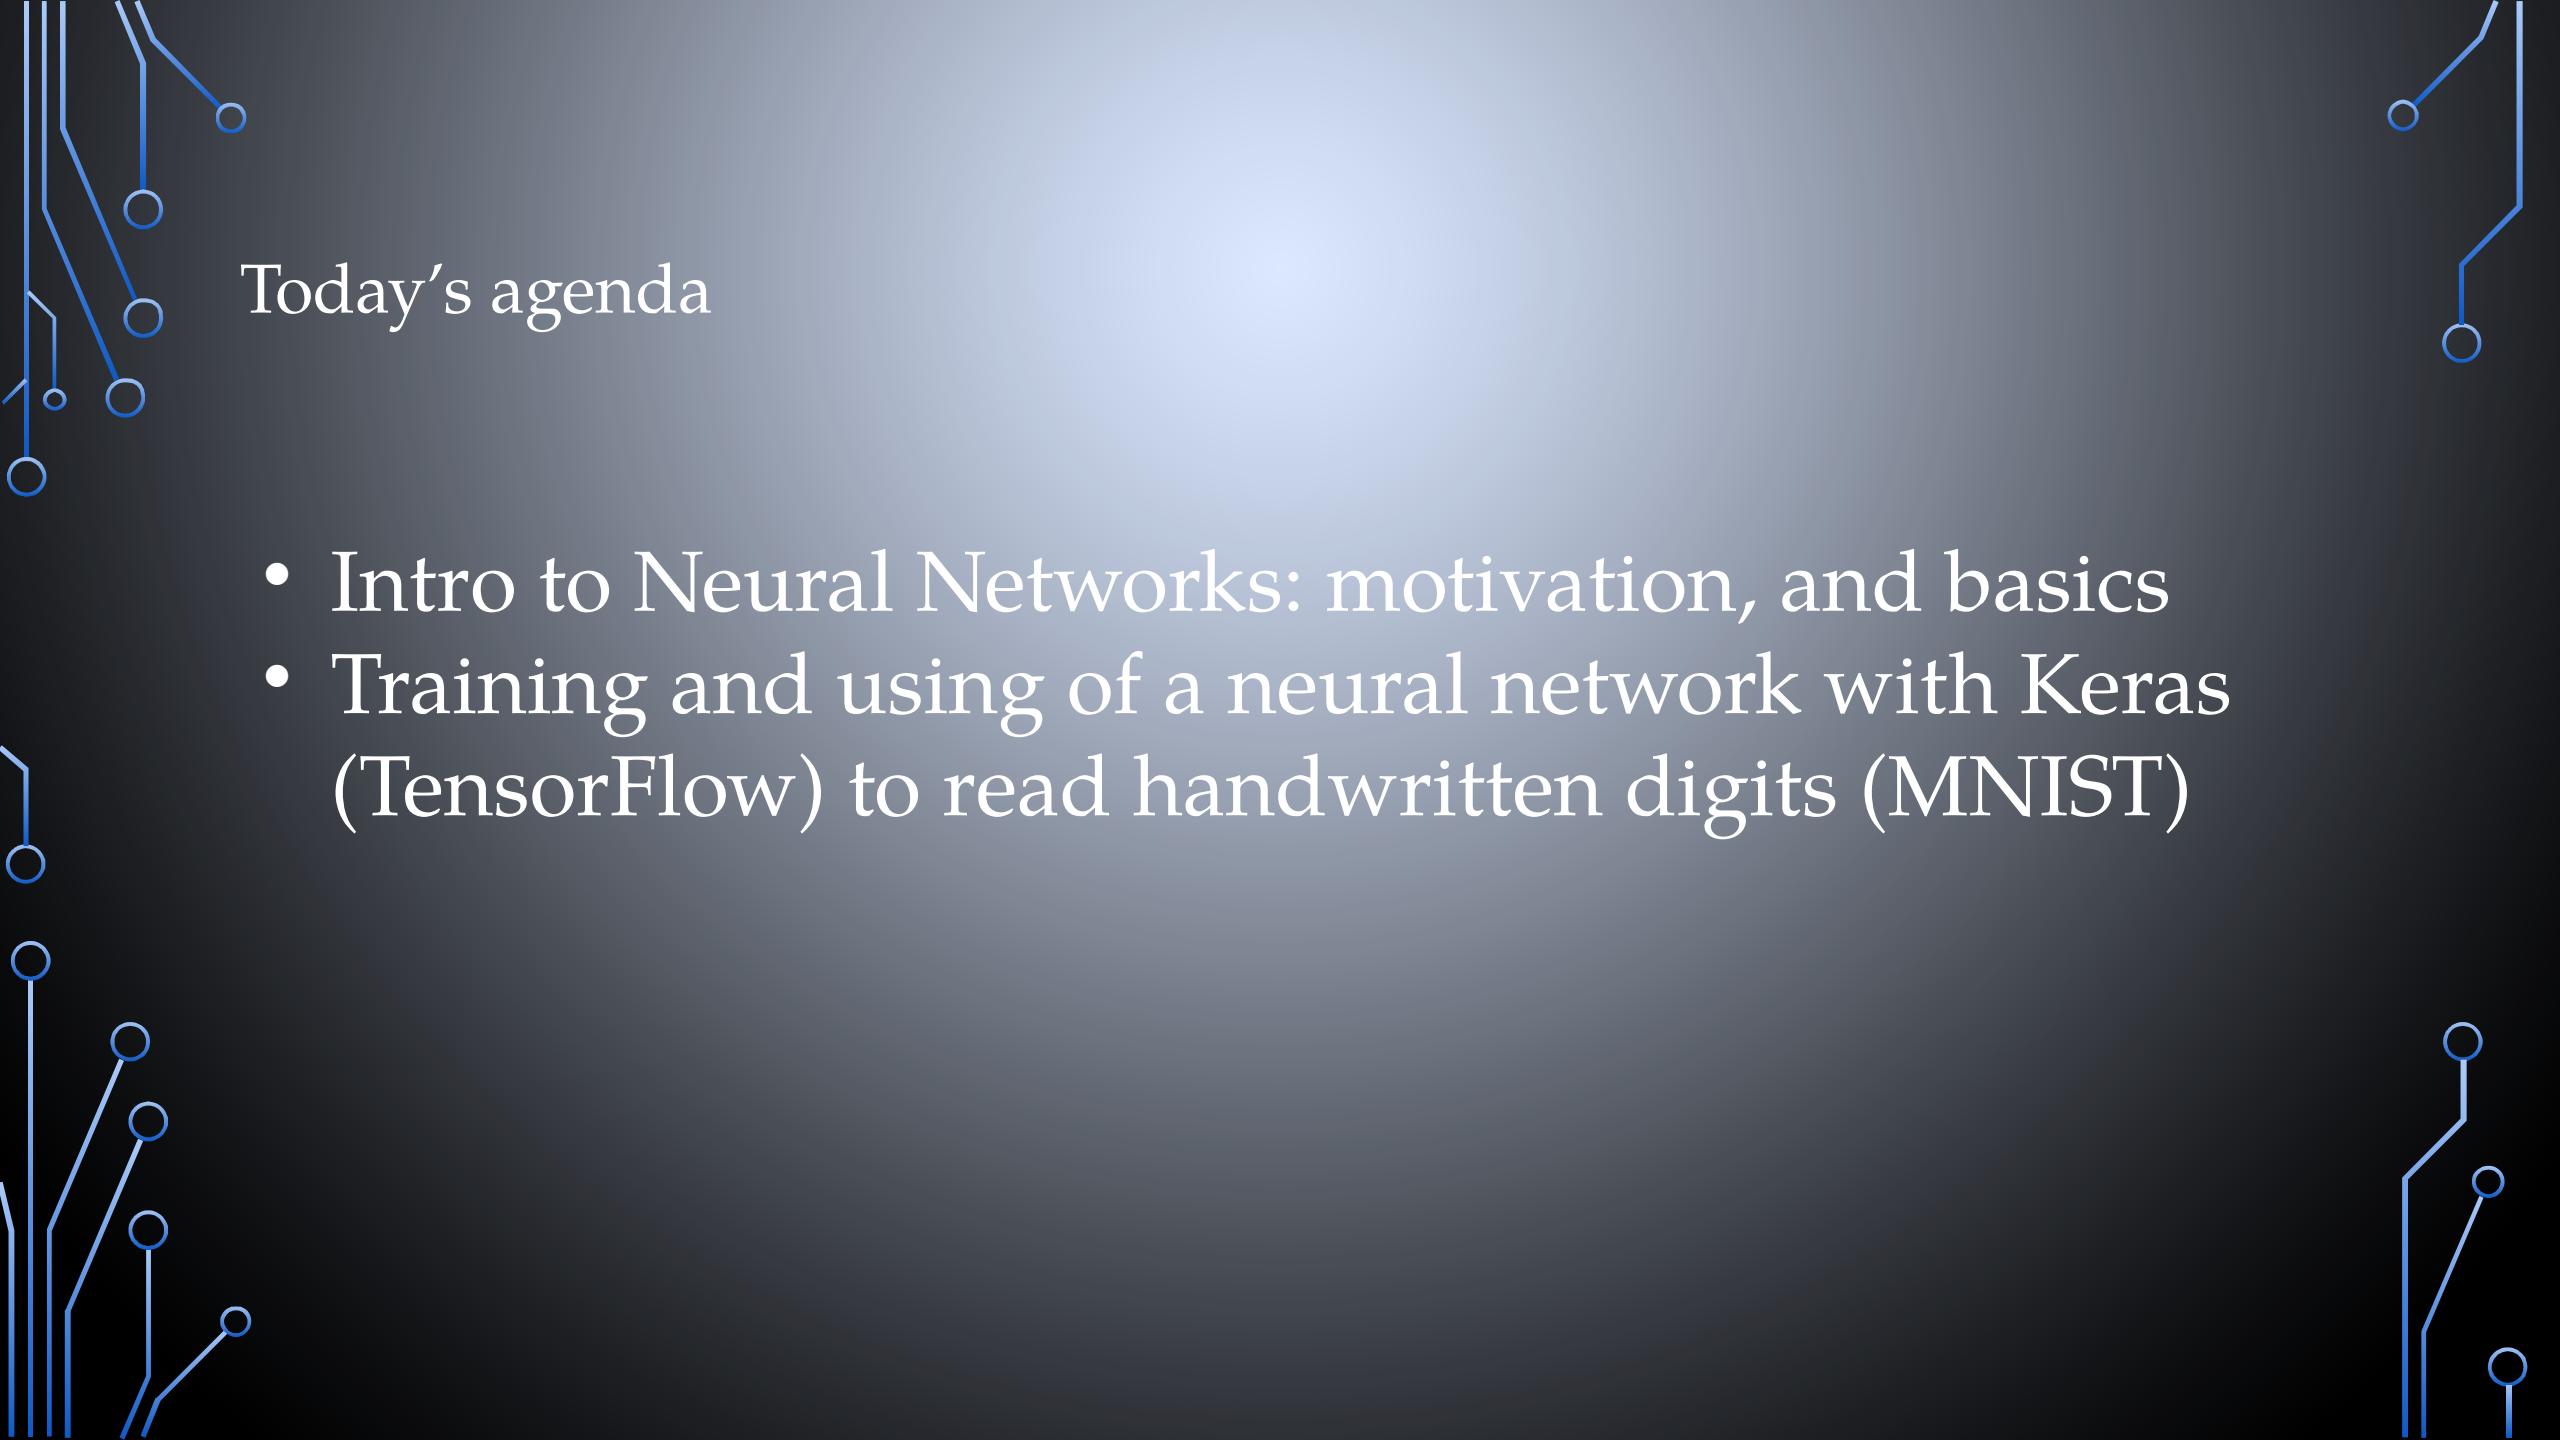




# LESSON 5: NEURAL NETWORKS

BY ALEXANDER MATHIS  
HARVARD UNIVERSITY

ORGANIZED BY BRUNO BOIVIN  
AND RAFAEL GONZALEZ CANO



## Today's agenda

- Intro to Neural Networks: motivation, and basics
- Training and using of a neural network with Keras (TensorFlow) to read handwritten digits (MNIST)

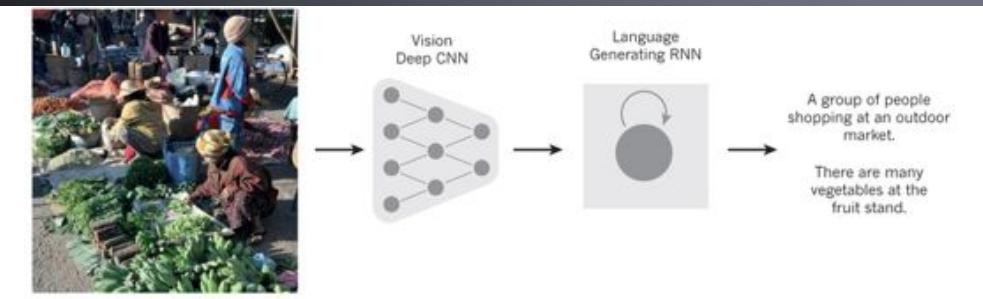
# Why neural networks?

- State-of-the art performance for many challenging tasks....

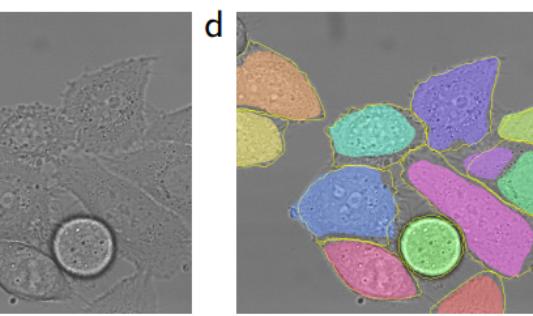


Object recognition,  
Alexnet 2012

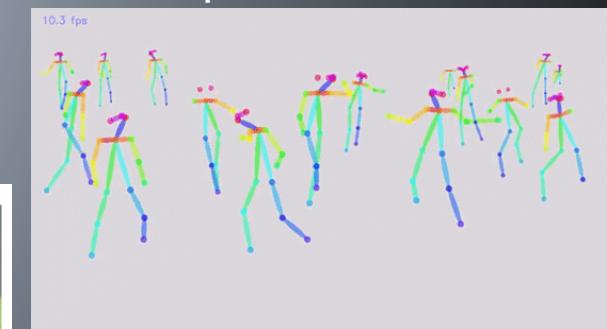
Speech recognition,  
Image >>> text,



Deep Learning,  
LeCun et al.  
Nature 2015



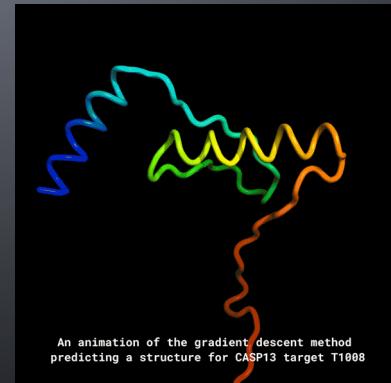
Cell tracking, U-Net



OpenPose



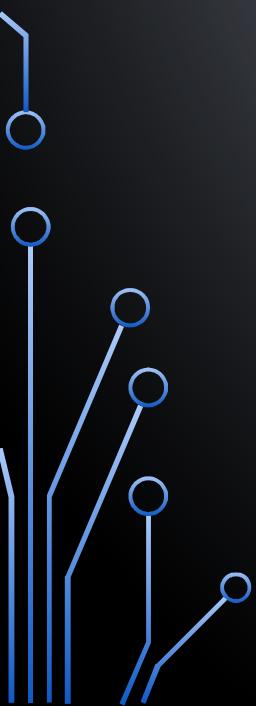
DeepLabCut





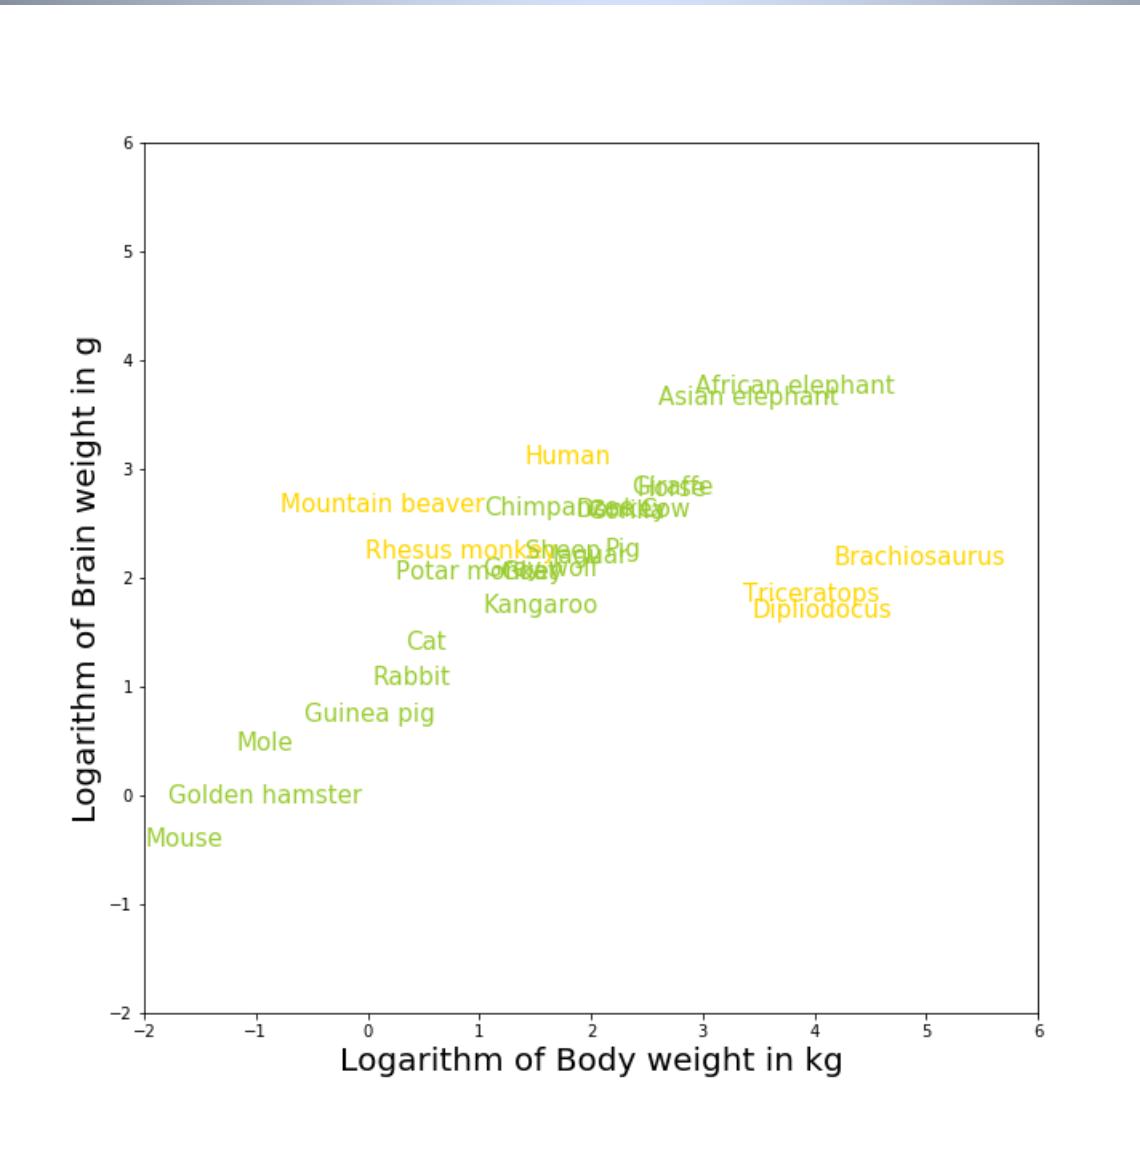
## What is learning?

Learning is the process of acquiring new, or modifying existing, knowledge, behaviors, skills, values, or preferences.  
*(wikipedia)*



For today... learning is curve fitting!

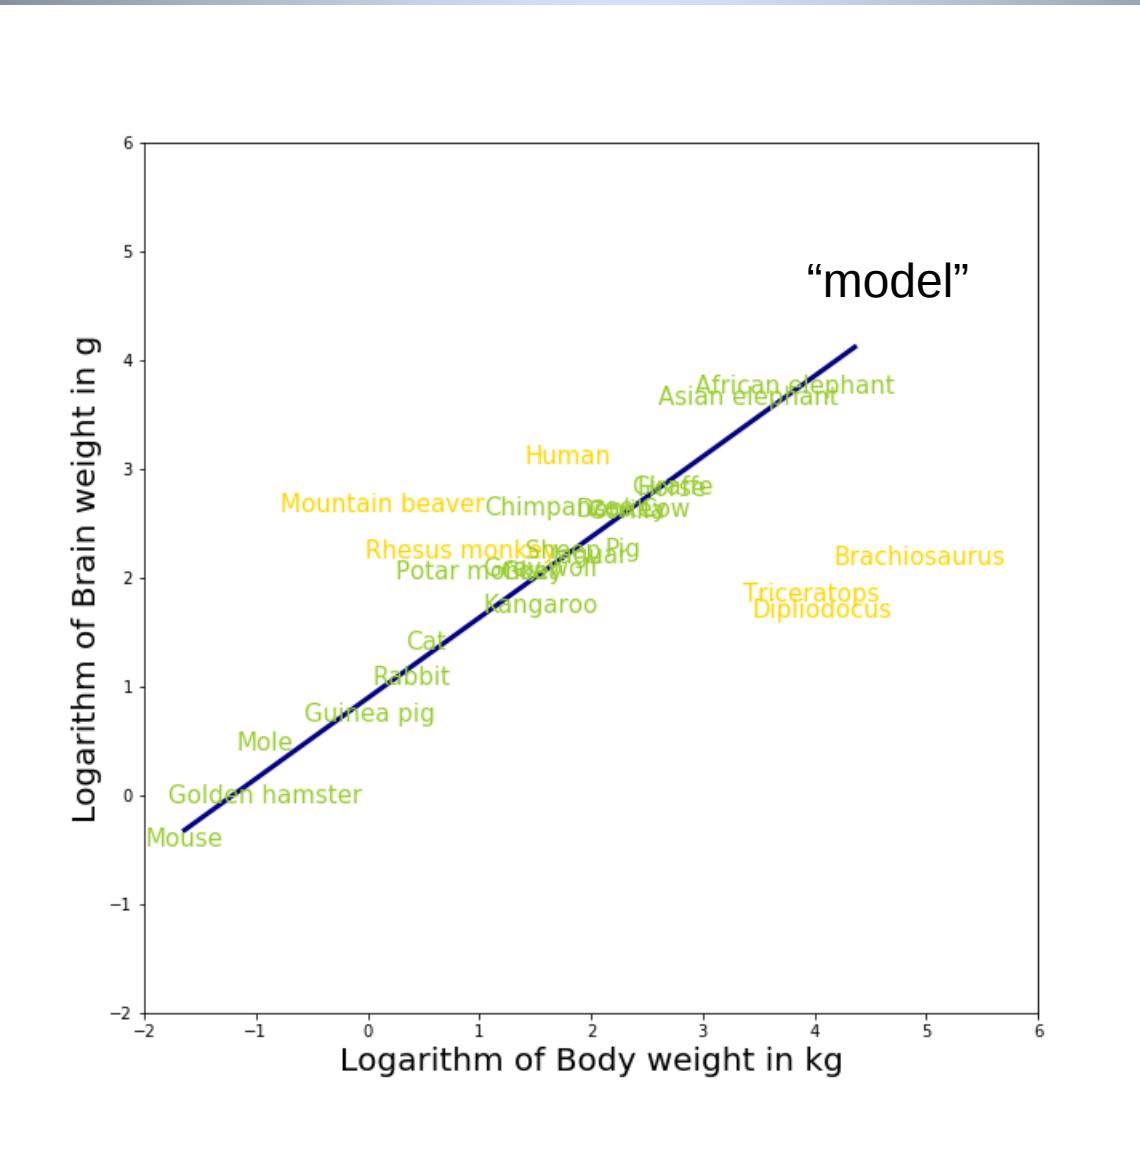
# Example 1 – scientific data



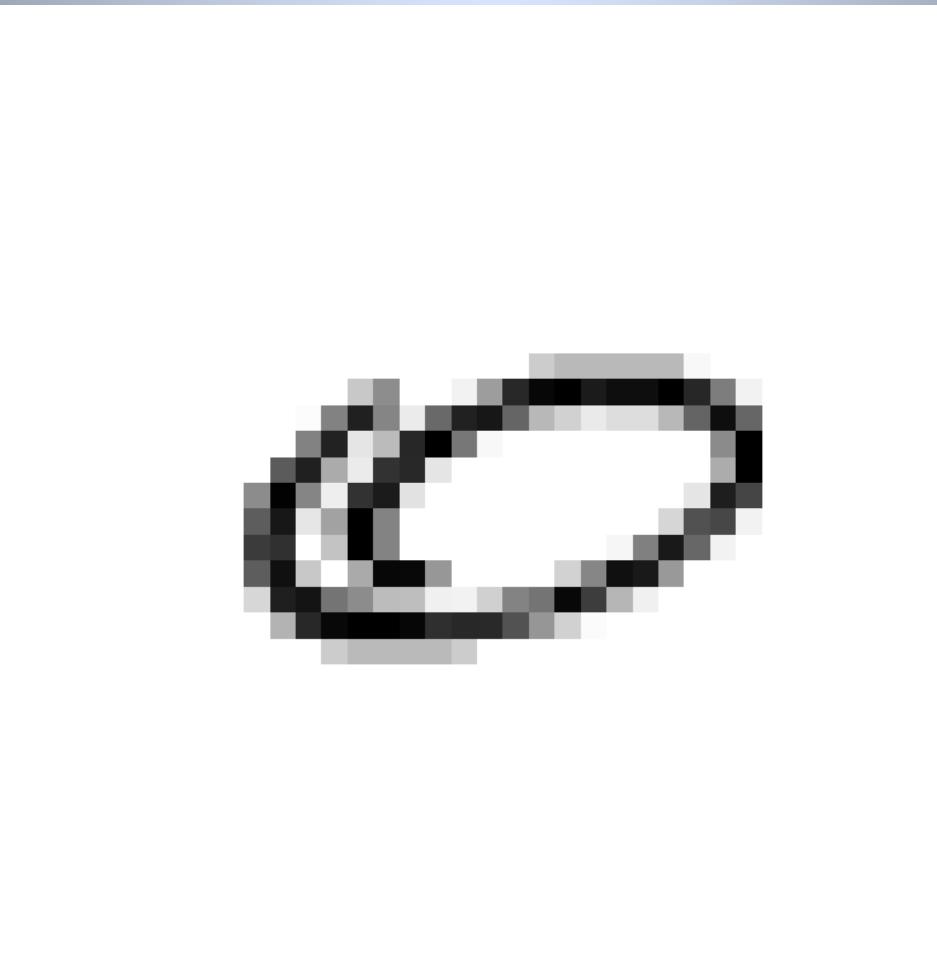
Average brain and body weights for 27 species of land animals.

Source: <http://mste.illinois.edu/malcz/DATA/BIOLOGY/Animals.html>

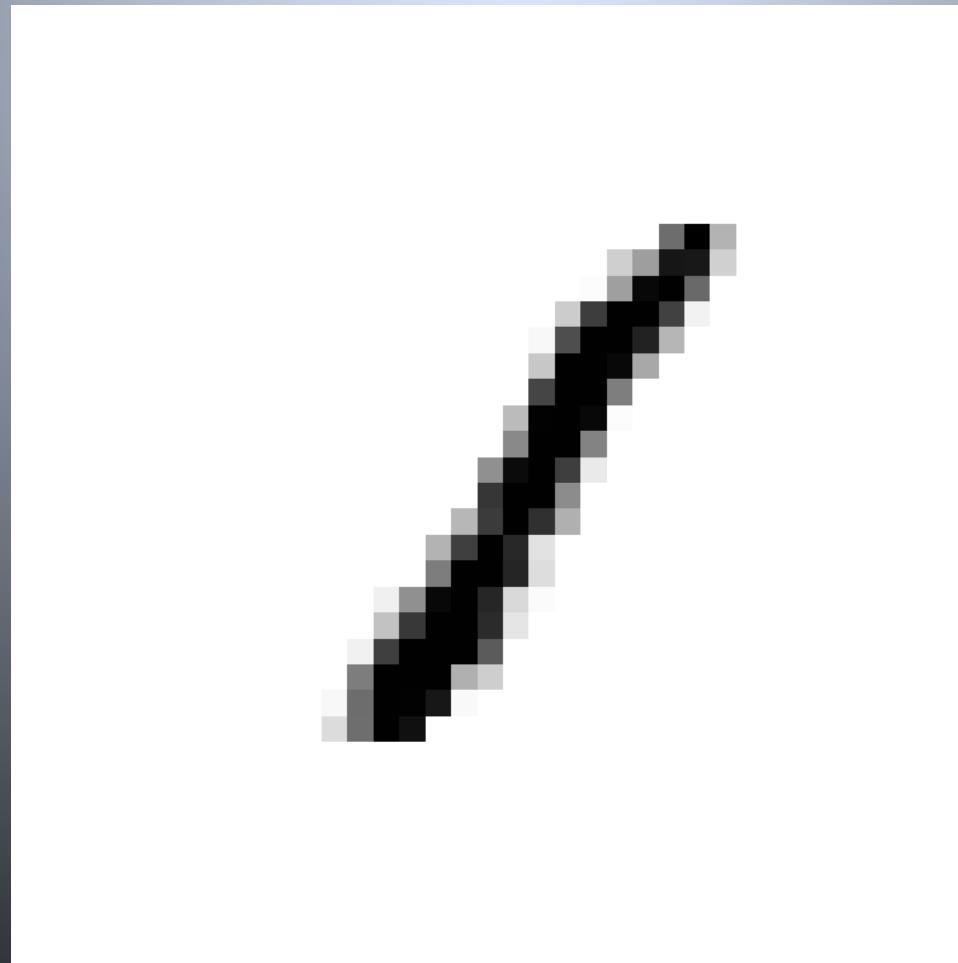
# Example 1 – scientific data



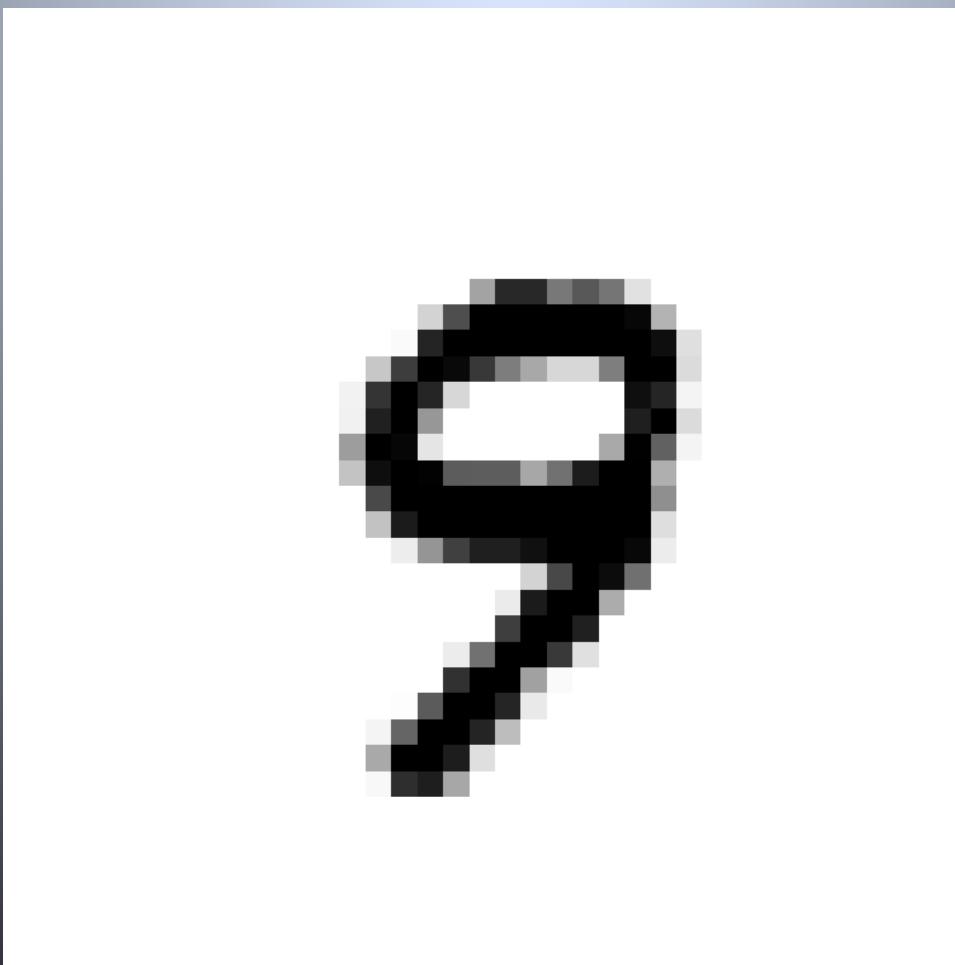
## Example 2 – perception



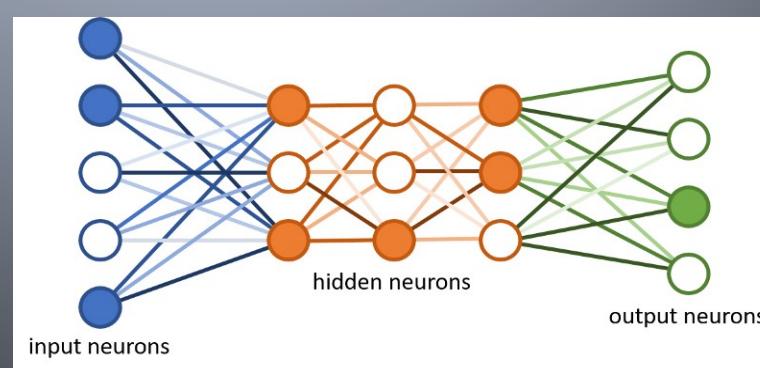
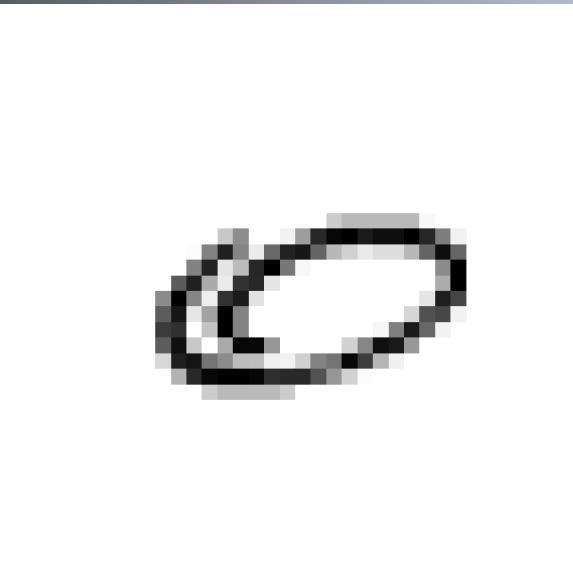
## Example 2 – perception



## Example 2 – perception



## Example 2 – perception



0?  
1?  
2?  
..

The MNIST database of handwritten digits – classical problem for reading handwritten digits  
<http://yann.lecun.com/exdb/mnist/>

Further resources....



The following slides are from  
Prof. Russ Salakhutdinov (Carnegie Mellon)

His full lecture slides can be found on the MLSS website!

- Machine Learning Summer School <http://mlss.tuebingen.mpg.de/2017/index.html>
- Deep Learning Summer School Materials  
<https://sites.google.com/site/deeplearningsummerschool2016/>

# ARTIFICIAL NEURON

- Neuron pre-activation (or input activation):

$$a(\mathbf{x}) = b + \sum_i w_i x_i = b + \mathbf{w}^\top \mathbf{x}$$

- Neuron output activation:

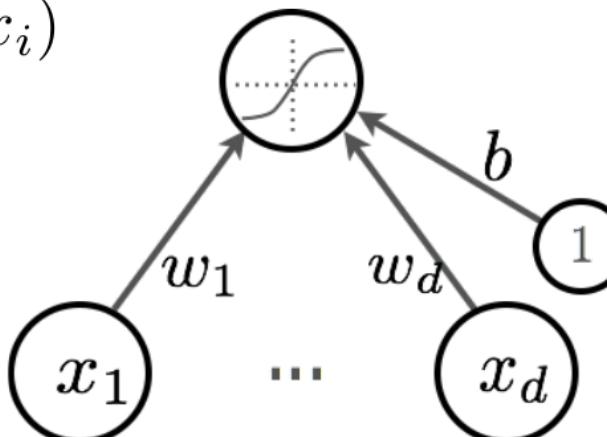
$$h(\mathbf{x}) = g(a(\mathbf{x})) = g(b + \sum_i w_i x_i)$$

where

$\mathbf{w}$  are the weights (parameters)

$b$  is the bias term

$g(\cdot)$  is called the activation function

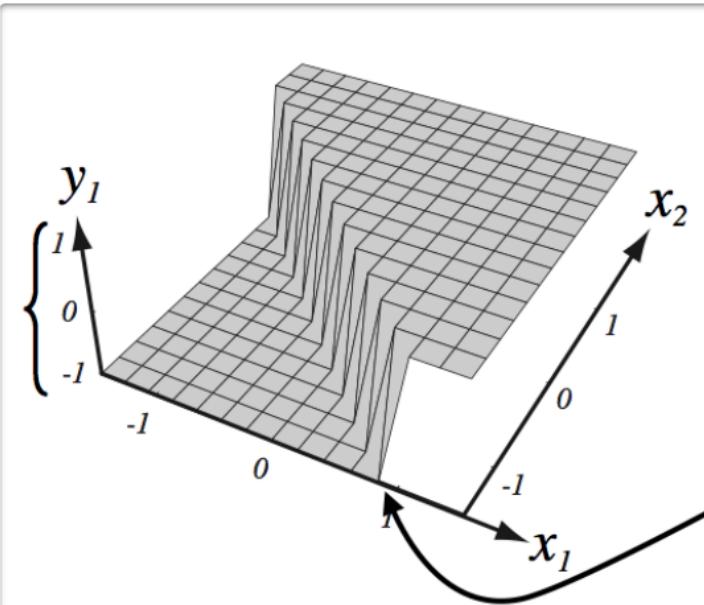


# ARTIFICIAL NEURON

- Output activation of the neuron:

$$h(\mathbf{x}) = g(a(\mathbf{x})) = g(b + \sum_i w_i x_i)$$

Range is determined by  $g(\cdot)$



(from Pascal Vincent's slides)

Bias only changes the position of the riff

# SINGLE HIDDEN LAYER NET

- Hidden layer pre-activation:

$$\mathbf{a}(\mathbf{x}) = \mathbf{b}^{(1)} + \mathbf{W}^{(1)}\mathbf{x}$$

$$(a(\mathbf{x})_i = b_i^{(1)} + \sum_j W_{i,j}^{(1)}x_j)$$

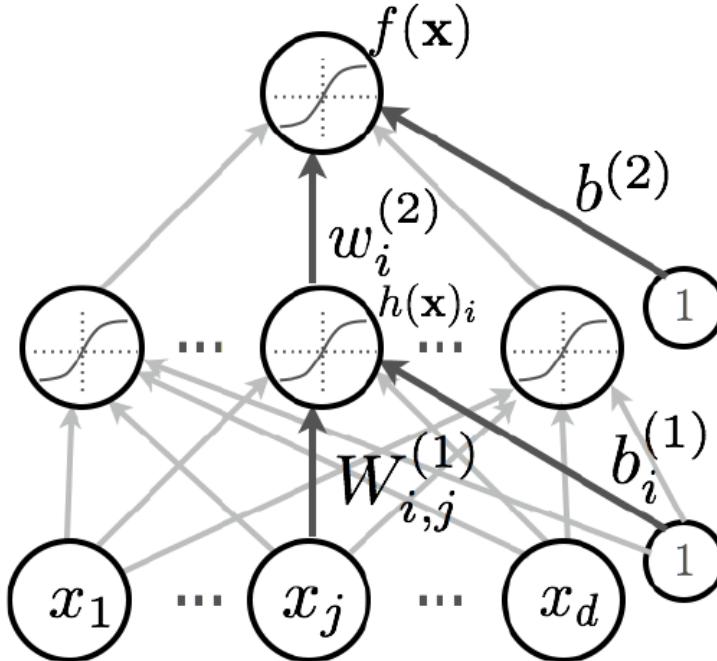
- Hidden layer activation:

$$\mathbf{h}(\mathbf{x}) = \mathbf{g}(\mathbf{a}(\mathbf{x}))$$

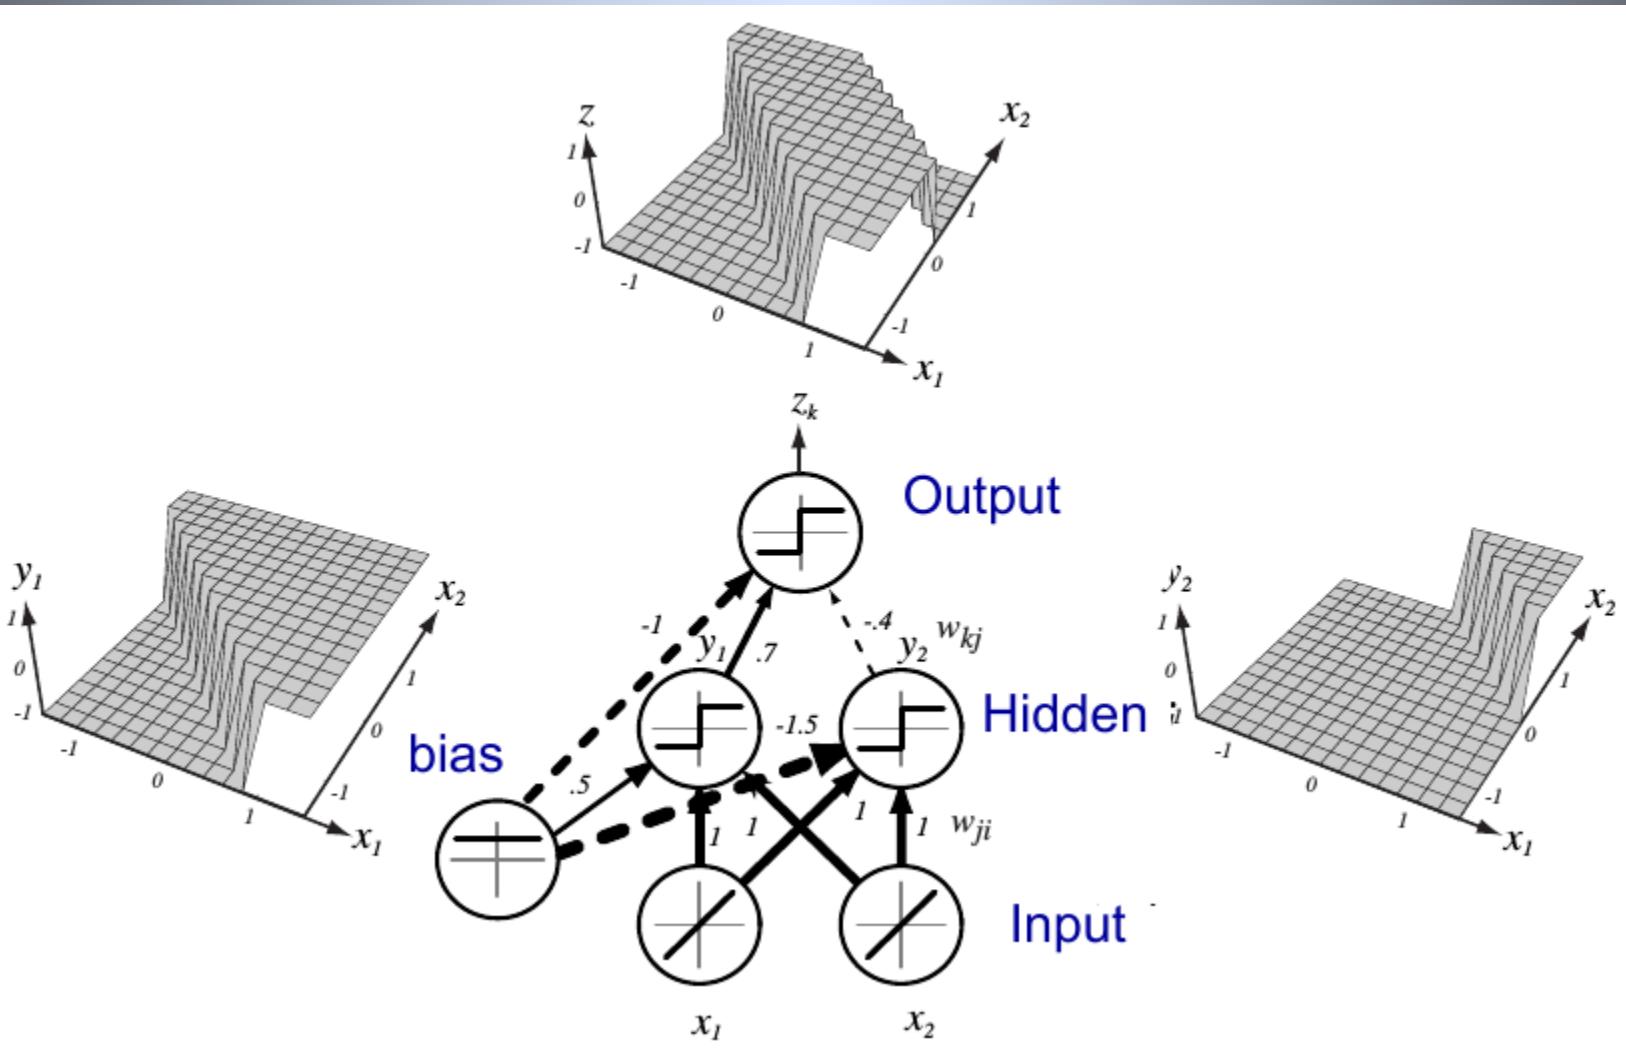
- Output layer activation:

$$f(\mathbf{x}) = o\left(b^{(2)} + \mathbf{w}^{(2)^\top} \mathbf{h}^{(1)} \mathbf{x}\right)$$

Output activation  
function

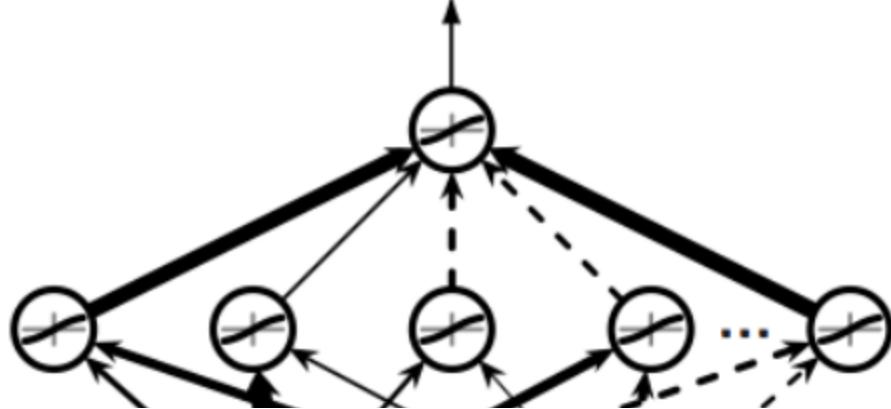


# CAPACITY



From Prof. Salakhutdinov

# CAPACITY



*Neural Networks*, Vol. 4, pp. 251-257, 1991  
Printed in the USA. All rights reserved.

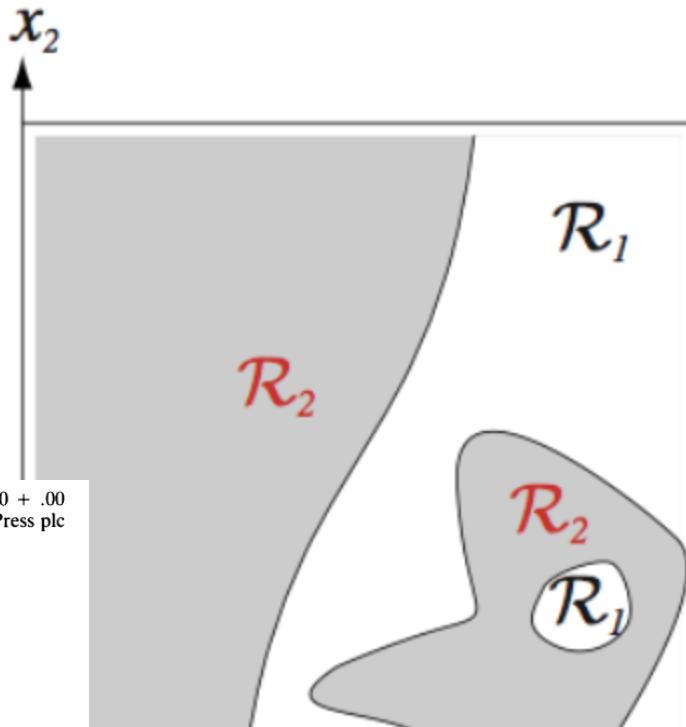
0893-6080/91 \$3.00 + .00  
Copyright © 1991 Pergamon Press plc

ORIGINAL CONTRIBUTION

## Approximation Capabilities of Multilayer Feedforward Networks

KURT HORNICK

**Theorem 1:** If  $\psi$  is unbounded and nonconstant, then  $\mathcal{H}_k(\psi)$  is dense in  $L^p(\mu)$  for all finite measures  $\mu$  on  $\mathbb{R}^k$ .



$$\mathcal{H}_k^{(n)}(\psi) = \left\{ h: \mathbb{R}^k \rightarrow \mathbb{R} \mid h(x) = \sum_{j=1}^n \beta_j \psi(a_j' x - \theta_j) \right\}$$

mutdinov

# MULTILAYER NEURAL NET

- Consider a network with L hidden layers.

- layer pre-activation for  $k > 0$

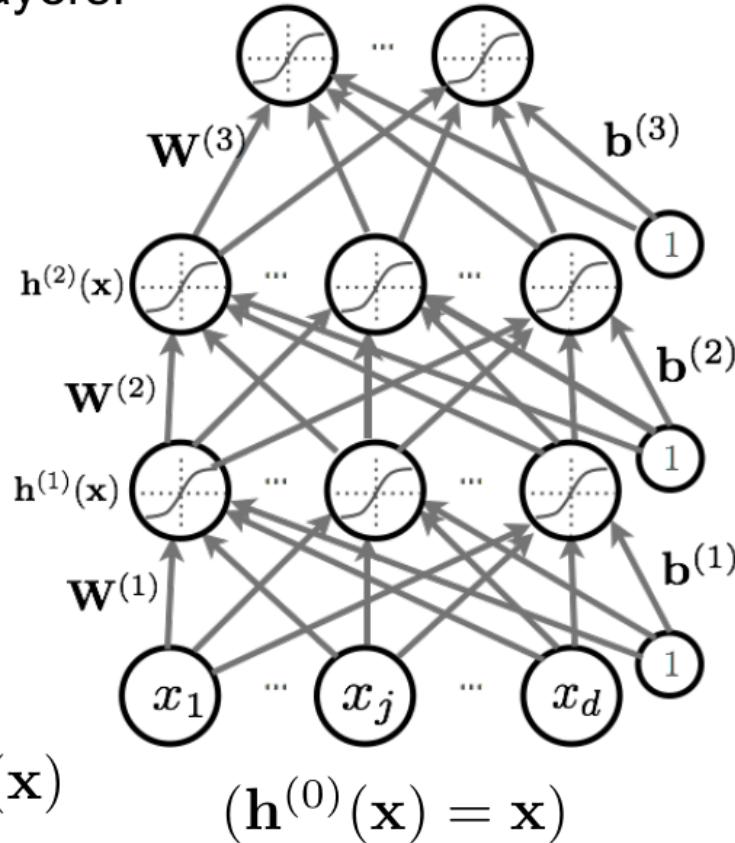
$$\mathbf{a}^{(k)}(\mathbf{x}) = \mathbf{b}^{(k)} + \mathbf{W}^{(k)} \mathbf{h}^{(k-1)}(\mathbf{x})$$

- hidden layer activation from 1 to L:

$$\mathbf{h}^{(k)}(\mathbf{x}) = \mathbf{g}(\mathbf{a}^{(k)}(\mathbf{x}))$$

- output layer activation ( $k=L+1$ ):

$$\mathbf{h}^{(L+1)}(\mathbf{x}) = \mathbf{o}(\mathbf{a}^{(L+1)}(\mathbf{x})) = \mathbf{f}(\mathbf{x})$$



# WELL BUT HOW DO WE LEARN?

- Empirical Risk Minimization:

$$\arg \min_{\theta} \frac{1}{T} \sum_t l(f(\mathbf{x}^{(t)}; \boldsymbol{\theta}), y^{(t)}) + \lambda \Omega(\boldsymbol{\theta})$$



Loss function                                      Regularizer

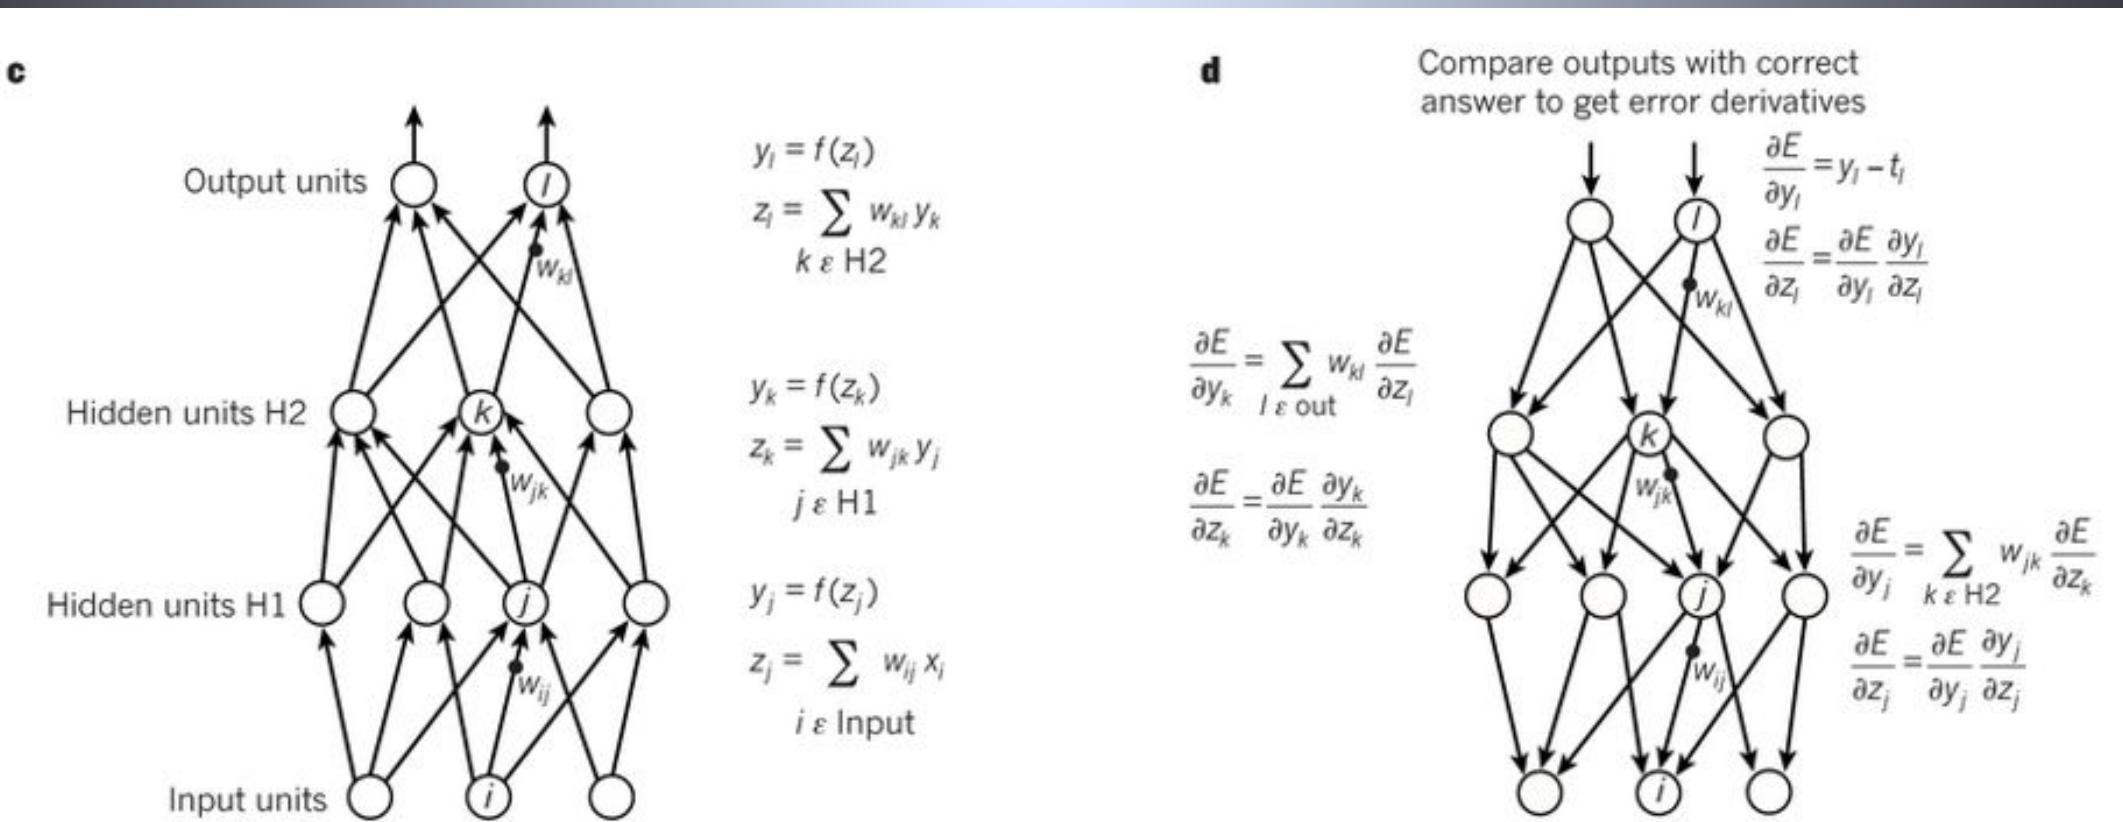
- To train a neural net, we need:

- **Loss function:**  $l(f(\mathbf{x}^{(t)}; \boldsymbol{\theta}), y^{(t)})$
- A procedure to **compute gradients**:  $\nabla_{\theta} l(f(\mathbf{x}^{(t)}; \boldsymbol{\theta}), y^{(t)})$
- **Regularizer and its gradient:**  $\Omega(\boldsymbol{\theta}), \nabla_{\theta} \Omega(\boldsymbol{\theta})$

# WELL BUT HOW DO WE LEARN?

- Perform updates after seeing each example:
  - Initialize:  $\theta \equiv \{\mathbf{W}^{(1)}, \mathbf{b}^{(1)}, \dots, \mathbf{W}^{(L+1)}, \mathbf{b}^{(L+1)}\}$
  - For  $t=1:T$ 
    - for each training example  $(\mathbf{x}^{(t)}, y^{(t)})$ 
$$\Delta = -\nabla_{\theta} l(f(\mathbf{x}^{(t)}; \theta), y^{(t)}) - \lambda \nabla_{\theta} \Omega(\theta)$$
$$\theta \leftarrow \theta + \alpha \Delta$$
- To train a neural net, we need:
  - **Loss function:**  $l(f(\mathbf{x}^{(t)}; \theta), y^{(t)})$
  - A procedure to **compute gradients:**  $\nabla_{\theta} l(f(\mathbf{x}^{(t)}; \theta), y^{(t)})$
  - **Regularizer and its gradient:**  $\Omega(\theta), \nabla_{\theta} \Omega(\theta)$

# BACKPROPAGATION (OF ERRORS) ALGORITHM

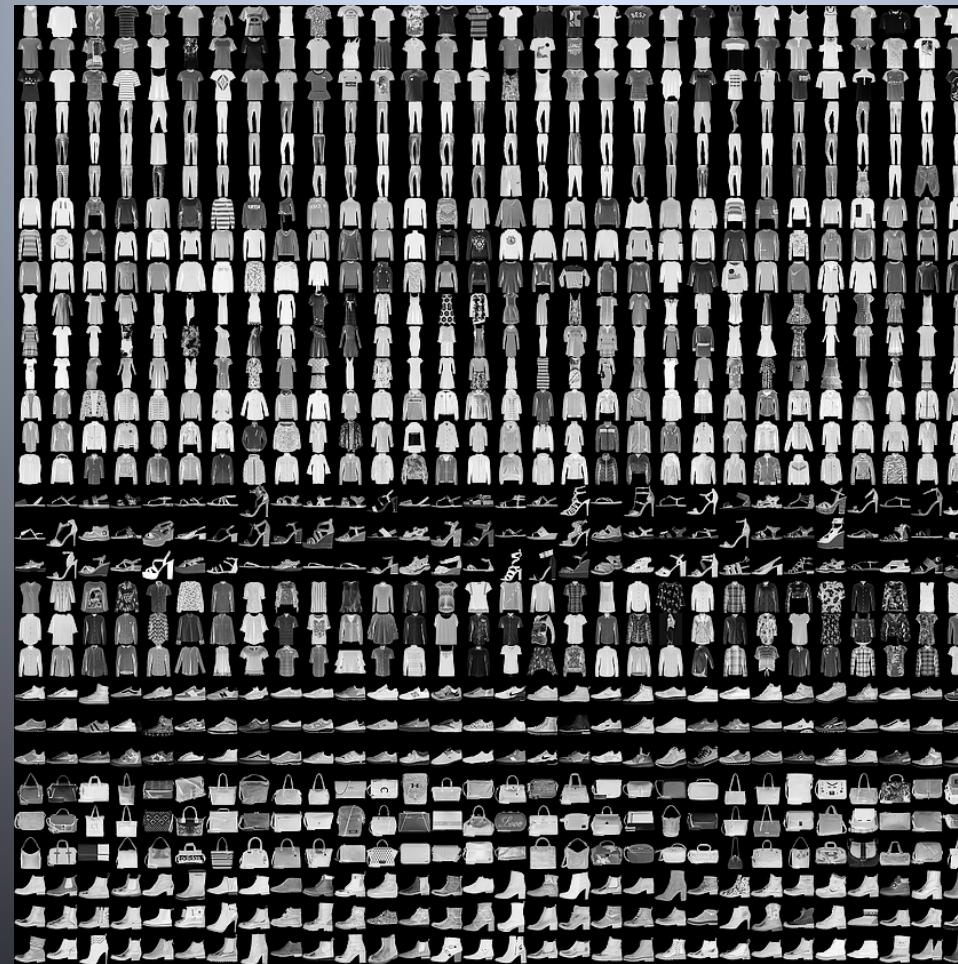




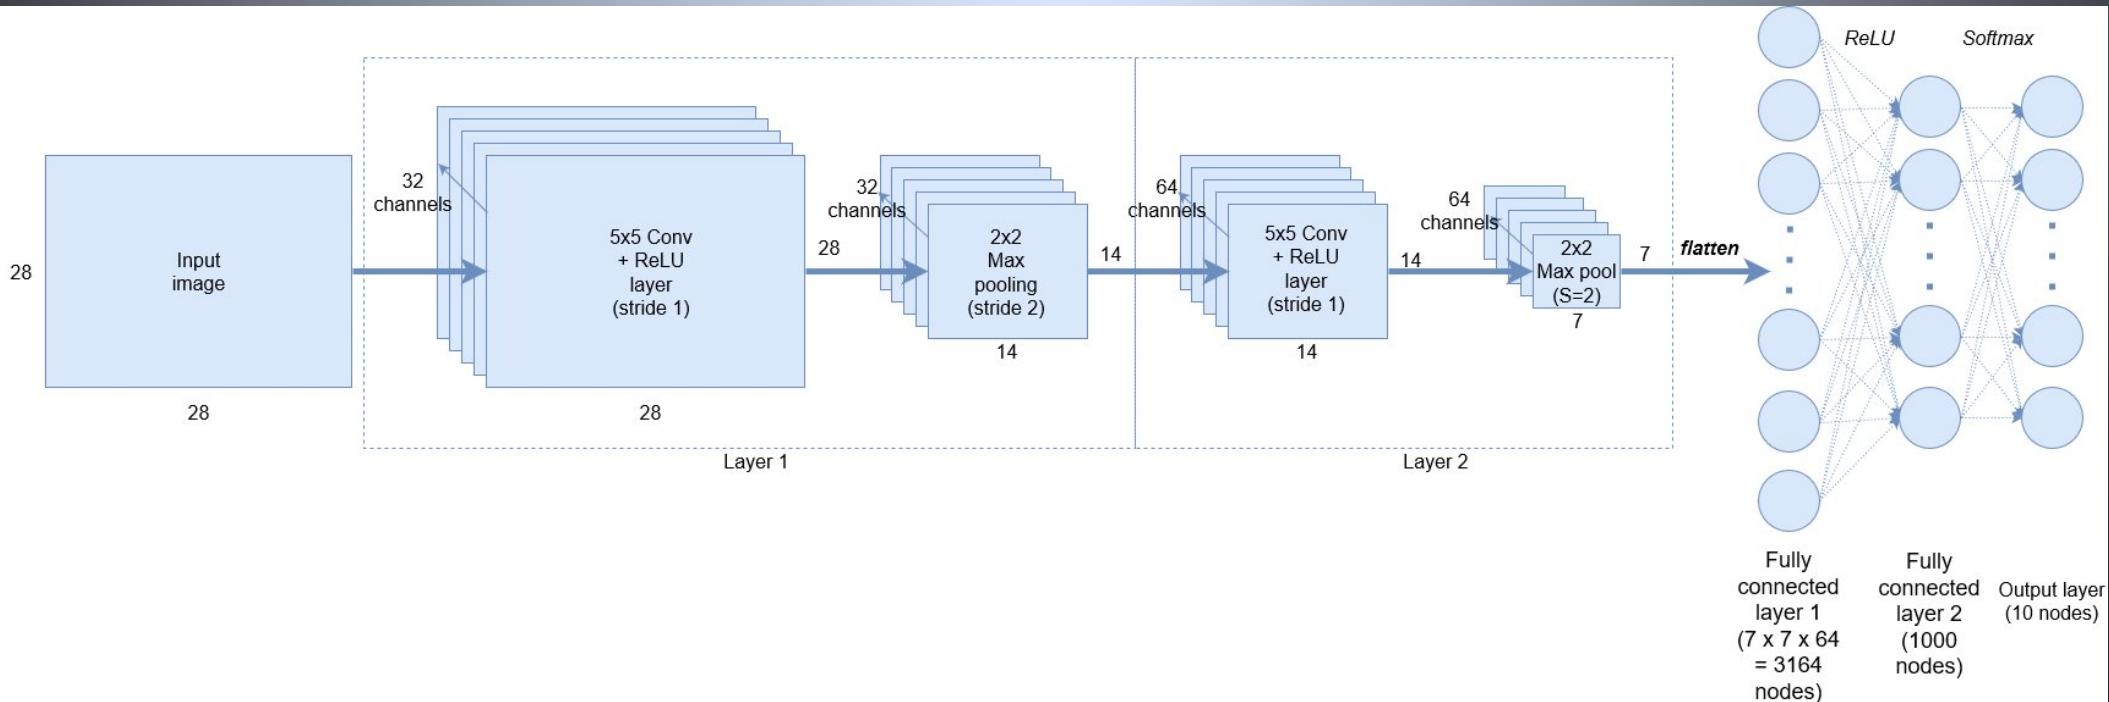
Fortunately,  
TensorFlow, Keras, Theano, PyTorch...  
minimize the loss for us!

Keras demo with MNIST and fashion MNIST....

<https://github.com/AlexEMG/Tutorials/blob/master/MNIST-Keras.ipynb>



# Convolutional Network architecture for Keras model with MNIST performance > 99%



In Keras this is as easy as:  
(see notebook for training etc.)

```
model = tf.keras.models.Sequential()  
model.add(tf.keras.layers.Conv2D(32, kernel_size=(5, 5), strides=(1, 1),  
                               activation='relu'))  
model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))  
model.add(tf.keras.layers.Conv2D(64, (5, 5), activation='relu'))  
model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))  
model.add(tf.keras.layers.Flatten())  
model.add(tf.keras.layers.Dense(1000, activation='relu'))  
model.add(tf.keras.layers.Dense(10, activation='softmax'))
```

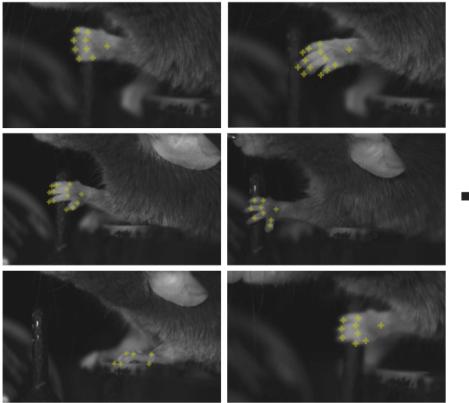
# Transfer learning...

## DeepLabCut Toolbox

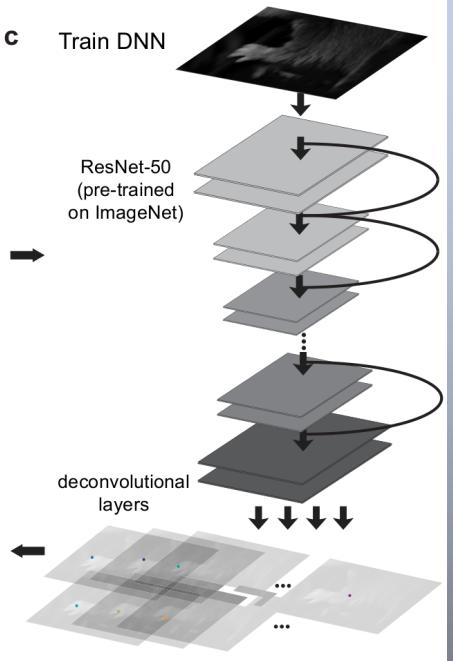
a Extract characteristic frames to label



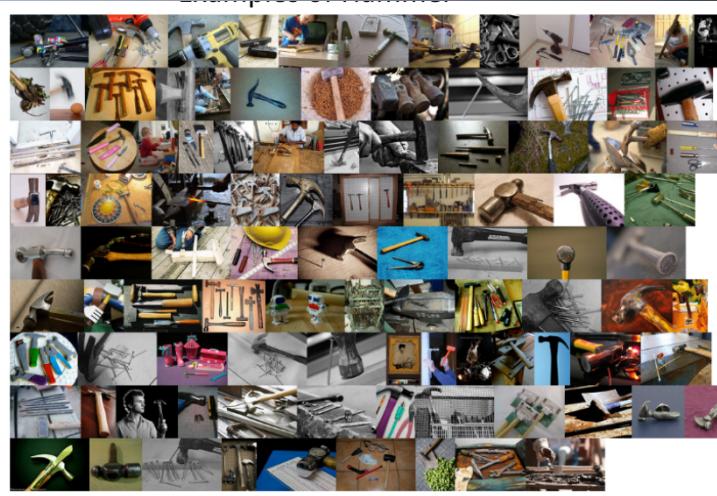
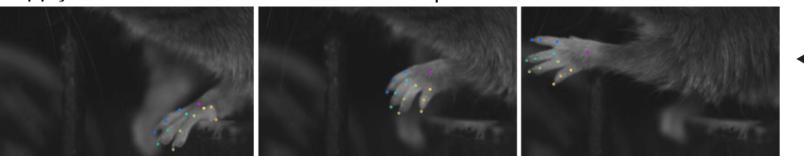
b Label features in frames



c Train DNN



d Apply to datasets: use trained network to predict labels



ImageNet (>1Mio images  
With 1,000 classes)

ResNet  
He et al. 2016

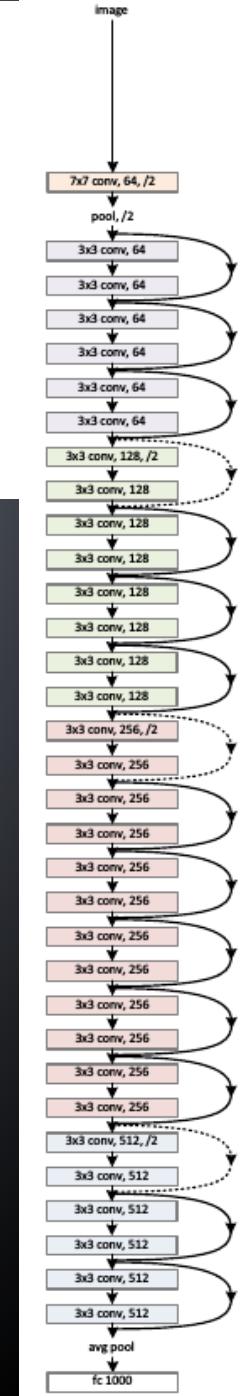
nature  
neuroscience

TECHNICAL REPORT

<https://doi.org/10.1038/s41593-018-0209-y>

## DeepLabCut: markerless pose estimation of user-defined body parts with deep learning

Alexander Mathis<sup>1,2</sup>, Pranav Mamidanna<sup>1</sup>, Kevin M. Cury<sup>3</sup>, Taiga Abe<sup>3</sup>, Venkatesh N. Murthy<sup>1,2</sup>, Mackenzie Weygandt Mathis<sup>1,4,8\*</sup> and Matthias Bethge<sup>1,5,6,7,8</sup>



A few links:

- Deep Learning textbook by Goodfellow, Bengio, Courville
- TensorFlow tutorial: <https://github.com/chiphuyen/stanford-tensorflow-tutorials>
- Awesome interpretability techniques for DeepNets  
<https://distill.pub/2018/building-blocks/>
- Linking neural nets to topology  
<http://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>  
<https://cs.stanford.edu/people/karpathy/convnetjs//demo/classify2d.html>
- Machine Learning Summer School  
<http://mlss.tuebingen.mpg.de/2017/index.html>
- Deep Learning Summer School Materials  
<https://sites.google.com/site/deeplearningsummerschool2016>



[info@neurodev.org](mailto:info@neurodev.org)  
[amathis@fas.harvard.edu](mailto:amathis@fas.harvard.edu)